

Algorithm and Data Structure Analysis (ADSA)

Lecture 5: Linear-Time Sorting Algorithms

Sorting So Far

- Insertion sort:
 - Easy to code
 - Fast on small inputs (less than ~50 elements)
 - Fast on nearly-sorted inputs
 - $O(n^2)$ worst case
 - $O(n^2)$ average (equally-likely inputs) case
 - $O(n^2)$ reverse-sorted case

Sorting So Far

- Merge sort:
 - Divide-and-conquer:
 - Split array in half
 - Recursively sort subarrays
 - Linear-time merge step
 - $O(n \lg n)$ worst case
 - Doesn't sort in place

Sorting So Far

- Heap sort:
 - Uses the very useful heap data structure
 - Complete binary tree
 - Heap property: parent key $>$ children's keys
 - $O(n \lg n)$ worst case
 - Sorts in place
 - Fair amount of shuffling memory around

Sorting So Far

- Quick sort:
 - Divide-and-conquer:
 - Partition array into two subarrays, recursively sort
 - All of first subarray < all of second subarray
 - No merge step needed!
 - $O(n \lg n)$ average case
 - Fast in practice
 - $O(n^2)$ worst case
 - Naïve implementation: worst case on sorted input
 - Address this with randomized quicksort

How Fast Can We Sort?

- We will provide a lower bound, then beat it
 - *How do you suppose we'll beat it?*
- First, an observation: all of the sorting algorithms so far are *comparison sorts*
 - The only operation used to gain ordering information about a sequence is the pairwise comparison of two elements
 - Theorem: all comparison sorts are $\Omega(n \lg n)$
 - A comparison sort must do $O(n)$ comparisons (*why?*)
 - What about the gap between $O(n)$ and $O(n \lg n)$

Decision Trees

- *Decision trees* provide an abstraction of comparison sorts
 - A decision tree represents the comparisons made by a comparison sort. Every thing else ignored
 - (Draw examples on board)
- *What do the leaves represent?*
- *How many leaves must there be?*

Decision Trees

- Decision trees can model comparison sorts.
For a given algorithm:
 - One tree for each n
 - Tree paths are all possible execution traces
 - *What's the longest path in a decision tree for insertion sort? For merge sort?*
- *What is the asymptotic height of any decision tree for sorting n elements?*
- Answer: $\Omega(n \lg n)$ (now let's prove it...)

Lower Bound For Comparison Sorting

- Thm: Any decision tree that sorts n elements has height $\Omega(n \lg n)$
- *What's the minimum # of leaves?*
- *What's the maximum # of leaves of a binary tree of height h ?*
- Clearly the minimum # of leaves is less than or equal to the maximum # of leaves

Lower Bound For Comparison Sorting

- So we have...

$$n! \leq 2^h$$

- Taking logarithms:

$$\lg(n!) \leq h$$

- Stirling's approximation tells us:

$$n! > \left(\frac{n}{e}\right)^n$$

- Thus: $h \geq \lg\left(\frac{n}{e}\right)^n$

David Luebke

Lower Bound For Comparison Sorting

- So we have

$$h \geq \lg \left(\frac{n}{e} \right)^n$$

$$= n \lg n - n \lg e$$

- $= \Omega(n \lg n)$
Thus the minimum height of a decision tree is $\Omega(n \lg n)$

Lower Bound For Comparison Sorts

- Thus the time to comparison sort n elements is $\Omega(n \lg n)$
- Corollary: Heapsort and Mergesort are asymptotically optimal comparison sorts
- But the name of this lecture is “Sorting in linear time”!
 - *How can we do better than $\Omega(n \lg n)$?*

Sorting In Linear Time

- Counting sort
 - No comparisons between elements!
 - *But...* depends on assumption about the numbers being sorted
 - We assume numbers are in the range $1..k$
 - The algorithm:
 - Input: $A[1..n]$, where $A[j] \in \{1, 2, 3, \dots, k\}$
 - Output: $B[1..n]$, sorted (notice: not sorting in place)
 - Also: Array $C[1..k]$ for auxiliary storage

Counting Sort

```
1  CountingSort(A, B, k)
2      for i=1 to k
3          C[i]= 0;
4      for j=1 to n
5          C[A[j]] += 1;
6      for i=2 to k
7          C[i] = C[i] + C[i-1];
8      for j=n downto 1
9          B[C[A[j]]] = A[j];
10         C[A[j]] -= 1;
```

Work through example: $A=\{4\ 1\ 3\ 4\ 3\}$, $k=4$

Counting Sort

```
1 CountingSort(A, B, k)
```

```
2   for i=1 to k
```

```
3     C[i] = 0;
```

```
4   for j=1 to n
```

```
5     C[A[j]] += 1;
```

```
6   for i=2 to k
```

```
7     C[i] = C[i] + C[i-1];
```

```
8   for j=n downto 1
```

```
9     B[C[A[j]]] = A[j];
```

```
10    C[A[j]] -= 1;
```

← Takes time $O(k)$

← Takes time $O(n)$

What will be the running time?

Counting Sort

- Total time: $O(n + k)$
 - Usually, $k = O(n)$
 - Thus counting sort runs in $O(n)$ time
- But sorting is $\Omega(n \lg n)$!
 - No contradiction--this is not a comparison sort (in fact, there are *no* comparisons at all!)
 - Notice that this algorithm is *stable*

Counting Sort

- Cool! *Why don't we always use counting sort?*
- Because it depends on range k of elements
- *Could we use counting sort to sort 32 bit integers? Why or why not?*
- Answer: no, k too large ($2^{32} = 4,294,967,296$)

Counting Sort

- *How did IBM get rich originally?*
- Answer: punched card readers for census tabulation in early 1900' s.
 - In particular, a *card sorter* that could sort cards into different bins
 - Each column can be punched in 12 places
 - Decimal digits use 10 places
 - Problem: only one column can be sorted on at a time

Radix Sort

- Intuitively, you might sort on the most significant digit, then the second msd, etc.
- Problem: lots of intermediate piles of cards (read: scratch arrays) to keep track of
- Key idea: sort the *least* significant digit first

RadixSort(A, d)

for i=1 to d

StableSort(A) on digit i

– Example: Fig 9.3

Radix Sort

- *Can we prove it will work?*
- Sketch of an inductive argument (induction on the number of passes):
 - Assume lower-order digits $\{j: j < i\}$ are sorted
 - Show that sorting next digit i leaves array correctly sorted
 - If two digits at position i are different, ordering numbers by that digit is correct (lower-order digits irrelevant)
 - If they are the same, numbers are already sorted on the lower-order digits. Since we use a stable sort, the numbers stay in the right order

Radix Sort

- *What sort will we use to sort on digits?*
- Counting sort is obvious choice:
 - Sort n numbers on digits that range from $1..k$
 - Time: $O(n + k)$
- Each pass over n numbers with d digits takes time $O(n+k)$, so total time $O(dn+dk)$
 - When d is constant and $k=O(n)$, takes $O(n)$ time
- *How many bits in a computer word?*

Radix Sort

- Problem: sort 1 million 64-bit numbers
 - Treat as four-digit radix 2^{16} numbers
 - Can sort in just four passes with radix sort!
- Compares well with typical $O(n \lg n)$ comparison sort
 - Requires approx $\lg n = 20$ operations per number being sorted
- *So why would we ever use anything but radix sort?*

Radix Sort

- In general, radix sort based on counting sort is
 - Fast
 - Asymptotically fast (i.e., $O(n)$)
 - Simple to code
 - A good choice
- To think about: *Can radix sort be used on floating-point numbers?*