

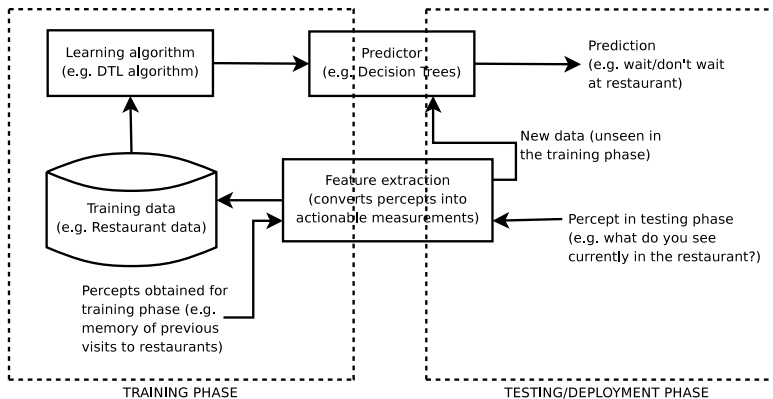
Assessing Performance, Preventing Overfitting

3007/7059 Artificial Intelligence

School of Computer Science
The University of Adelaide

Basic pipeline of learning algorithms

There are 2 major stages in a learning algorithm:

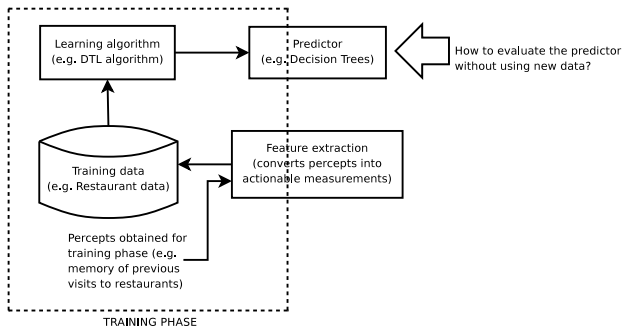


Note that the training and testing data are obtained at **different stages in time**.

A good learning system makes accurate predictions...

In most cases we would like to evaluate the predictor by how well it predicts the label for new samples/data, i.e. its performance in the **testing phase**. This implies, however, that we would have to **deploy** the system first before evaluating it— A very risky strategy!

Ideally we want to evaluate before deployment, but how do we test prediction capability on **unseen** data before deployment?

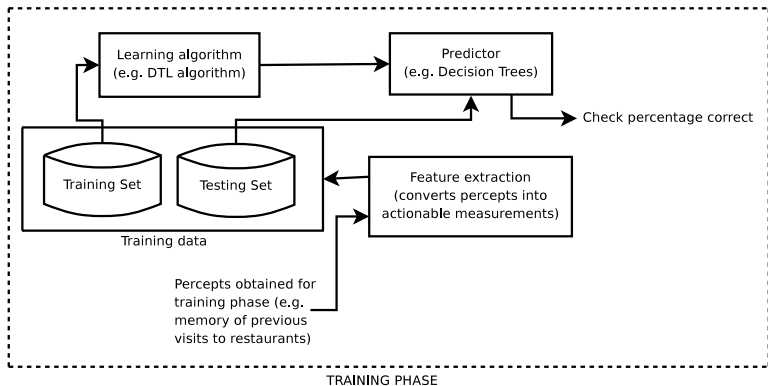


Methodology for Assessing Performance

What we can do is as follows:

1. Collect a large set of training data, i.e. measurements and their desired labels.
2. **Randomly** divide it into two **disjoint (non-overlapping)** subsets: the **training set** and the **testing set**.
3. Apply the learning algorithm (e.g. decision trees) on the training set, producing a predictor/classifier.
4. Measure the percentage of samples in the testing set that are correctly labelled by the predictor/classifier.
5. **Repeat** steps 2–4 for **different sizes** of training sets.


Methodology for Assessing Performance (cont.)



Methodology for Assessing Performance (cont.)

The results can be plotted on a graph called the **learning curve**, i.e., a plot of the accuracy (% of correct prediction) versus size of training set.

Example: Applying the DTL algorithm on the “restaurant” problem with 100 randomly generated samples:



Notice that the correct prediction rate increases with the size of the training set, a sign that the learning algorithm is picking up the pattern in the data.

Peeking or Cheating

Peeking or Cheating: Allowing the learning algorithm to access or “see” the testing set before performance evaluation.

This is usually (inadvertently) done in 2 ways:

1. Allowing the training and testing sets to overlap. An extreme case is to report performance based on prediction results on the training set itself, i.e. training set = testing set!
2. A more subtle way is to tweak the learning algorithm based on its performance on the testing set.


Always avoid peeking or cheating by ensuring absolute separation between the training and testing set.

Overfitting

Overfitting: Learning meaningless regularity or patterns in the data.

Overfitting is detrimental to the **generalisation capability** (i.e. the prediction performance) of a predictor/classifier.

Recall the 1D regression problem: The line which captures best the pattern of the data will have the highest generalisation capability.



Recall Occam's Razor which states that we should prefer the simplest hypothesis that explains the data.

Example: Overfitting in the restaurant problem


Data can have irrelevant attributes/features.

Lets add a new attribute (name of restaurant) to the restaurant data:

	Attributes										
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>Name</i>
<i>X</i> ₁	<i>T</i>	<i>F</i>	<i>F</i>	<i>T</i>	<i>Some</i>	<i>\$\$\$</i>	<i>F</i>	<i>T</i>	<i>French</i>	<i>0–10</i>	<i>CBistrot</i>
<i>X</i> ₂	<i>T</i>	<i>F</i>	<i>F</i>	<i>T</i>	<i>Full</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Thai</i>	<i>30–60</i>	<i>CBlossom</i>
<i>X</i> ₃	<i>F</i>	<i>T</i>	<i>F</i>	<i>F</i>	<i>Some</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Burger</i>	<i>0–10</i>	<i>HJacks</i>
<i>X</i> ₄	<i>T</i>	<i>F</i>	<i>T</i>	<i>T</i>	<i>Full</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Thai</i>	<i>10–30</i>	<i>LGrass</i>
<i>X</i> ₅	<i>T</i>	<i>F</i>	<i>T</i>	<i>F</i>	<i>Full</i>	<i>\$\$\$</i>	<i>F</i>	<i>T</i>	<i>French</i>	<i>>60</i>	<i>LGuillotine</i>
<i>X</i> ₆	<i>F</i>	<i>T</i>	<i>F</i>	<i>T</i>	<i>Some</i>	<i>\$\$</i>	<i>T</i>	<i>T</i>	<i>Italian</i>	<i>0–10</i>	<i>CPrimo</i>
<i>X</i> ₇	<i>F</i>	<i>T</i>	<i>F</i>	<i>F</i>	<i>None</i>	<i>\$</i>	<i>T</i>	<i>F</i>	<i>Burger</i>	<i>0–10</i>	<i>McD</i>
<i>X</i> ₈	<i>F</i>	<i>F</i>	<i>F</i>	<i>T</i>	<i>Some</i>	<i>\$\$</i>	<i>T</i>	<i>T</i>	<i>Thai</i>	<i>0–10</i>	<i>CThai</i>
<i>X</i> ₉	<i>F</i>	<i>T</i>	<i>T</i>	<i>F</i>	<i>Full</i>	<i>\$</i>	<i>T</i>	<i>F</i>	<i>Burger</i>	<i>>60</i>	<i>SWay</i>
<i>X</i> ₁₀	<i>T</i>	<i>T</i>	<i>T</i>	<i>T</i>	<i>Full</i>	<i>\$\$\$</i>	<i>F</i>	<i>T</i>	<i>Italian</i>	<i>10–30</i>	<i>RAmalfi</i>
<i>X</i> ₁₁	<i>F</i>	<i>F</i>	<i>F</i>	<i>F</i>	<i>None</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Thai</i>	<i>0–10</i>	<i>SiamR</i>
<i>X</i> ₁₂	<i>T</i>	<i>T</i>	<i>T</i>	<i>T</i>	<i>Full</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Burger</i>	<i>30–60</i>	<i>RRooster</i>

Example: Overfitting in the restaurant problem (cont.)

As long as no two samples have the same value for some attribute, the DTL algorithm will **always** find a **consistent** hypothesis!




A predictor or classifier that exhibits overfitting is unlikely to be useful.

DTL on Iris data


Trying to predict one of two types of irises based on measurements of the sepal length and width.

Each sample \mathbf{x}_i contains two values (sepal length, sepal width) and a class label y_i (*setosa* or *virginica*). The training set can be plotted in 2D:



DTL on Iris data (cont.)

Intuitively the data is composed of two distinct areas with a smooth curve separating the two types of irises:



DTL on Iris data (cont.)

In contrast to the Restaurant data, the Iris data attributes is **continuous**, i.e. the value of each attribute (sepal length, sepal width) can vary continuously within a range. In tabular form

	Attributes		<i>Type</i>
	<i>Sepal length</i>	<i>Sepal width</i>	
x_1	4.3	3.2	<i>Setosa</i>
x_2	4.4	2.9	<i>Setosa</i>
x_3	6.2	3.4	<i>Virginica</i>
x_4	4.5	2.3	<i>Setosa</i>
x_5	6.1	2.6	<i>Virginica</i>
x_6	6.1	3.0	<i>Virginica</i>
x_7	4.4	3.2	<i>Setosa</i>
x_8	6.2	2.8	<i>Virginica</i>
\vdots	\vdots	\vdots	\vdots

Intuitively, each attribute corresponds to a **dimension** of the continuous valued data.

DTL for continuous data

In contrast to DTL for discrete data, in DTL for continuous data we never run-out of attributes since we can always find a value to split along a particular dimension:

- ▶ function $\text{DTL}(\text{examples}, \text{default})$ returns a decision tree dtree
 - ▶ if examples is empty then return default
 - ▶ else if all examples have the same labels then return a single-node tree containing examples .
 - ▶ else
 - ▶ $(\text{bestdim}, \text{bestsplit}) \leftarrow \text{CHOOSE-SPLIT}(\text{examples})$
 - ▶ $\text{dtree} \leftarrow$ a new decision tree with root test $(\text{bestdim}, \text{bestsplit})$
 - ▶ $\text{left} \leftarrow$ elements of examples with value **less** than bestsplit at dimension bestdim
 - ▶ $\text{leftchild} \leftarrow \text{DTL}(\text{left}, \text{MODE}(\text{left}))$
 - ▶ Attach leftchild as a subtree to dtree
 - ▶ $\text{right} \leftarrow$ elements of examples with value **more** than bestsplit at dimension bestdim
 - ▶ $\text{rightchild} \leftarrow \text{DTL}(\text{right}, \text{MODE}(\text{right}))$
 - ▶ Attach rightchild as a subtree to dtree

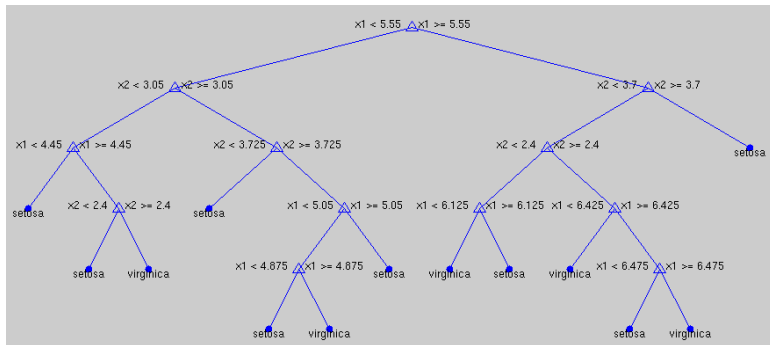
DTL for continuous data (cont.)

The subroutine CHOOSE-SPLIT basically finds the best dimension and value along which to split the examples according to the principle of maximising information gain:

- ▶ function CHOOSE-SPLIT(*examples*) returns (*bestdim*, *bestsplit*)
 - ▶ $bestdim \leftarrow NULL, bestsplit \leftarrow NULL, bestgain \leftarrow 0$
 - ▶ for each dimension d_i of *examples*
 - ▶ search for the value s_i along d_i to split *examples* which gives the highest information gain. Store this gain as $bestgain_i$.
 - ▶ if $bestgain_i > bestgain$, then $bestgain \leftarrow bestgain_i$, $bestdim \leftarrow d_i, bestsplit \leftarrow s_i$.


DTL for continuous data (cont.)

Running DTL on the Iris data produces the following Decision Tree:



Overfitting the Iris data

The predictions of the Decision Tree can be displayed as **decision areas** in the space of the data:



It can be seen that DTL is able to learn a Decision Tree that fits the data perfectly. However due to overfitting the Decision Tree will likely produce errors in the testing phase!

Decision Tree Pruning

A technique to prevent overfitting in DTL.

Two general strategies:

1. **Post-pruning:** This requires a separate testing set. First grow the tree fully and remove leaf nodes one-by-one if the prediction accuracy of the tree on the testing set improves. Stop as soon as the prediction accuracy starts to deteriorate.
2. **Pre-pruning:** Try to limit the growth of the tree by stopping prematurely. The following criteria can be used to decide when to stop splitting at a particular node:
 - ▶ Stop splitting as soon as the **number of data remaining** at the node is less than a pre-determined number.
OR
 - ▶ Stop splitting as soon as the **information content** at the node is less than a pre-determined minimum value.

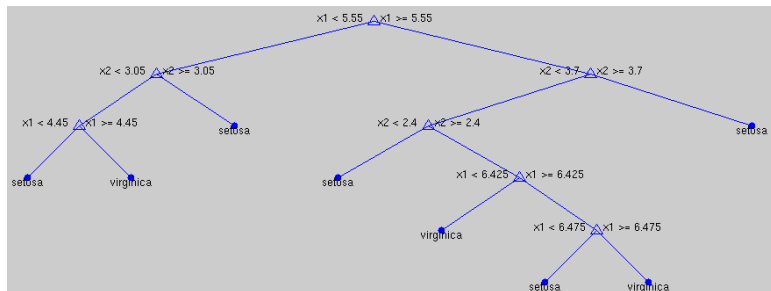
Pre-Pruning: Minimum number of data per node

Essentially we are just **changing the stopping criterion** of DTL such that it doesn't try to perfectly split the data:

- ▶ function $\text{DTL}(\text{examples}, \text{default}, \text{minsplit})$ returns a decision tree *dtree*
 - ▶ if *examples* is empty then return *default*
 - ▶ else if $\text{size}(\text{examples}) < \text{minsplit}$ then return a single-node tree containing *examples*.
 - ▶ else
 - ▶ $(\text{bestdim}, \text{bestsplit}) \leftarrow \text{CHOOSE-SPLIT}(\text{examples})$
 - ▶ *dtree* \leftarrow a new decision tree with root test $(\text{bestdim}, \text{bestsplit})$
 - ▶ *left* \leftarrow elements of *examples* with value less than *bestsplit* at dimension *bestdim*
 - ▶ *leftchild* $\leftarrow \text{DTL}(\text{left}, \text{MODE}(\text{left}), \text{minsplit})$
 - ▶ Attach *leftchild* as a subtree to *dtree*
 - ▶ *right* \leftarrow elements of *examples* with value more than *bestsplit* at dimension *bestdim*
 - ▶ *rightchild* $\leftarrow \text{DTL}(\text{right}, \text{MODE}(\text{right}), \text{minsplit})$
 - ▶ Attach *rightchild* as a subtree to *dtree*


Pre-Pruning: Minimum number of data per node (cont.)

For example, in the Iris data we impose $minsplitleft = 4$. The corresponding Decision Tree looks like:



Pre-Pruning: Minimum number of data per node (cont.)

The decision areas now look like:



Notice that some training examples are classified incorrectly by the pruned tree. However this Decision Tree would **generalise** better to new unseen data.


Pre-Pruning: Minimum information content per node

Another way is to impose a minimum information content at each node before allowing splitting:

- ▶ function $\text{DTL}(\text{examples}, \text{default}, \text{mininfo})$ returns a decision tree dtree
 - ▶ if examples is empty then return default
 - ▶ else if $I(\text{examples}) < \text{mininfo}$ then return a single-node tree containing examples .
 - ▶ else
 - ▶ $(\text{bestdim}, \text{bestsplit}) \leftarrow \text{CHOOSE-SPLIT}(\text{examples})$
 - ▶ $\text{dtree} \leftarrow$ a new decision tree with root test $(\text{bestdim}, \text{bestsplit})$
 - ▶ $\text{left} \leftarrow$ elements of examples with value less than bestsplit at dimension bestdim
 - ▶ $\text{leftchild} \leftarrow \text{DTL}(\text{left}, \text{MODE}(\text{left}), \text{mininfo})$
 - ▶ Attach leftchild as a subtree to dtree
 - ▶ $\text{right} \leftarrow$ elements of examples with value more than bestsplit at dimension bestdim
 - ▶ $\text{rightchild} \leftarrow \text{DTL}(\text{right}, \text{MODE}(\text{right}), \text{mininfo})$
 - ▶ Attach rightchild as a subtree to dtree

Pre-Pruning: Minimum information content per node (cont.)


For example, in the Iris data, we impose the minimum information content of 0.3 bits. The corresponding Decision Tree looks like:



You can verify that we stopped splitting as soon as the information content is less than 0.3 bits.

Pre-Pruning: Minimum information content per node (cont.)

The decision areas now look like:



Ensemble Learning

Ensemble Learning: To combine the outputs of an ensemble of hypotheses (predictors/classifiers) with the aim of improving prediction accuracy.

Consider an ensemble of M hypotheses h_1, h_2, \dots, h_M where we combine their predictions using a simple majority vote:

$$h_{overall}(x) = majority\{h_1(x), h_2(x), \dots, h_M(x)\}$$

Ensemble Learning (cont.)

For example, we have 5 Decision Trees h_1, h_2, \dots, h_5 trained with different pruning methods from the Iris data.

Given new data x each Decision Tree will make a prediction on which type of Iris it belongs to, e.g.

$h_1(x)=setosa, h_2(x)=virginica, h_3(x)=setosa, h_4(x)=setosa, h_5(x)=virginica$

The overall prediction on which type of Iris is simply

majority{*setosa, virginica, setosa, setosa, virginica*}

which amounts to *setosa*.

The question is, does this help in making more accurate predictions?

Ensemble Learning (cont.)

For the overall hypothesis to misclassify, it is clear that at least $(M - 1)/2 + 1$ ensemble members have to misclassify. From the previous example, at least $(5 - 1)/2 + 1 = 3$ Decision Trees have to misclassify before the overall prediction becomes wrong.

Assuming that each member hypothesis has a probability of error of p , and that the errors are independent, the overall probability of error is then

$$error = \sum_{i=(M-1)/2+1}^M \binom{M}{i} p^i (1-p)^{M-i}$$

where $\binom{M}{i}$ is the binomial coefficient

$$\binom{M}{i} = \frac{M!}{i!(M-i)!}$$

i.e. the number of ways to choose i items out of M items.

Ensemble Learning (cont.)

For example, if $M = 5$ and $p = 0.1$, the probability of error for the classifier ensemble is

$$\sum_{i=3}^5 \binom{5}{i} 0.1^i 0.9^{5-i} = \binom{5}{3} 0.1^3 0.9^2 + \binom{5}{4} 0.1^4 0.9^1 + \binom{5}{5} 0.1^5 0.9^0$$

which is less than 0.01!

Therefore, by combining a set of predictions from different hypotheses, we can actually reduce the overall probability of error.

Summary

Always separate examples/data into a training set and a testing set to facilitate the assessment of predictors/classifiers before deployment. To prevent cheating/peeking ensure that the training and testing sets are disjoint.

Overfitting is a serious problem in learning algorithms that can negatively impact the prediction accuracy of predictors/classifiers.

Decision Tree Pruning is a general method to minimise the effects of overfitting in Decision Tree learning.

Ensemble learning (aggregating/combining the outputs of several different predictors) is also an effective method to improve the prediction accuracy of a AI system.