


Domain Name System (DNS)

RFC 1034/1035

- The Internet identifies hosts by their IP address (e.g. 129.127.12.6)
- People tend to prefer names (e.g. www.cs.adelaide.edu.au)
- There are **lots** of Internet hosts. A centralised database wouldn't scale. The DNS is a **distributed** database.
 - organisations can change host names and IP addresses within their domain without informing a central authority
 - an organisation will typically have a **name server**
- A DNS server provides **name resolution** = conversion from a domain name to an IP address
 - a name server is a process listening on UDP/TCP port 53 for requests
 - when detected, the name is resolved, and a reply is sent
- Lots of applications are DNS clients, including web browsers.

DNS hierarchical database

- **Root name server**
 - (13 worldwide)
 - able to resolve all queries or identify another **intermediate** name server
- **Top-level domain (TLD) servers:**
 - responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
 - Network Solutions maintains servers for com TLD, Educause for edu TLD
- **Authoritative DNS servers:**
 - organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
 - can be maintained by organization or service provider
- **Local name server** - handles local DNS requests. Must know at least one root server. Caches resolved addresses.




DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:


- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?



DNS records

RR format: (name, value, type, ttl)

DNS: distributed db storing resource records (RR)

- Type=A
 - **name** is hostname
 - **value** is IP address
- Type=NS
 - **name** is domain (e.g. foo.com)
 - **value** is hostname of authoritative name server for this domain
- Type=CNAME
 - **name** is alias name for some “canonical” (the real) name
www.ibm.com is really
servereast.backup2.ibm.com
 - **value** is canonical name
- Type=MX
 - **value** is name of mailserver associated with **name**

Inserting records into DNS

- Example: just created startup “Network Utopia”
- Register name networkutopia.com at a **registrar** (e.g., Network Solutions)
 - Need to provide registrar with names and IP addresses of your authoritative name server (primary and secondary)
 - Registrar inserts two RRs into the com TLD server:

```
(networkutopia.com,  
  dns1.networkutopia.com, NS)  
(dns1.networkutopia.com, 212.212.212.1,  
  A)
```
- Put in authoritative server Type A record for `www.networkutopia.com` and Type MX record for `networkutopia.com`
- **How do people get the IP address of your Web site?**

Attacking DNS

DDoS attacks

- Bombard root servers with traffic
 - Not successful to date
 - Traffic Filtering
 - Local DNS servers cache IPs of TLD servers, allowing root server bypass
- Bombard TLD servers
 - Potentially more dangerous

Redirect attacks



- ❖ Man-in-middle
 - Intercept queries
- ❖ DNS poisoning
 - Send bogus replies to DNS server, which caches

Exploit DNS for DDoS

- ❖ Send queries with spoofed source address: target IP
- ❖ Requires amplification


Models of Interaction

- Client - Server
 - central storage of information is always on server
 - distinction between client which receives service and server which provides service
 - note that it is possible for a host to act as both a client and as a server in different interactions.
 - Web, e-mail, FTP
- Peer to Peer
 - distributed storage of information
 - no clear distinction between clients and servers. Hosts share typically equal control of processing and data
 - Peers dynamically join and leave
 - Bit Torrent



Pure P2P architecture

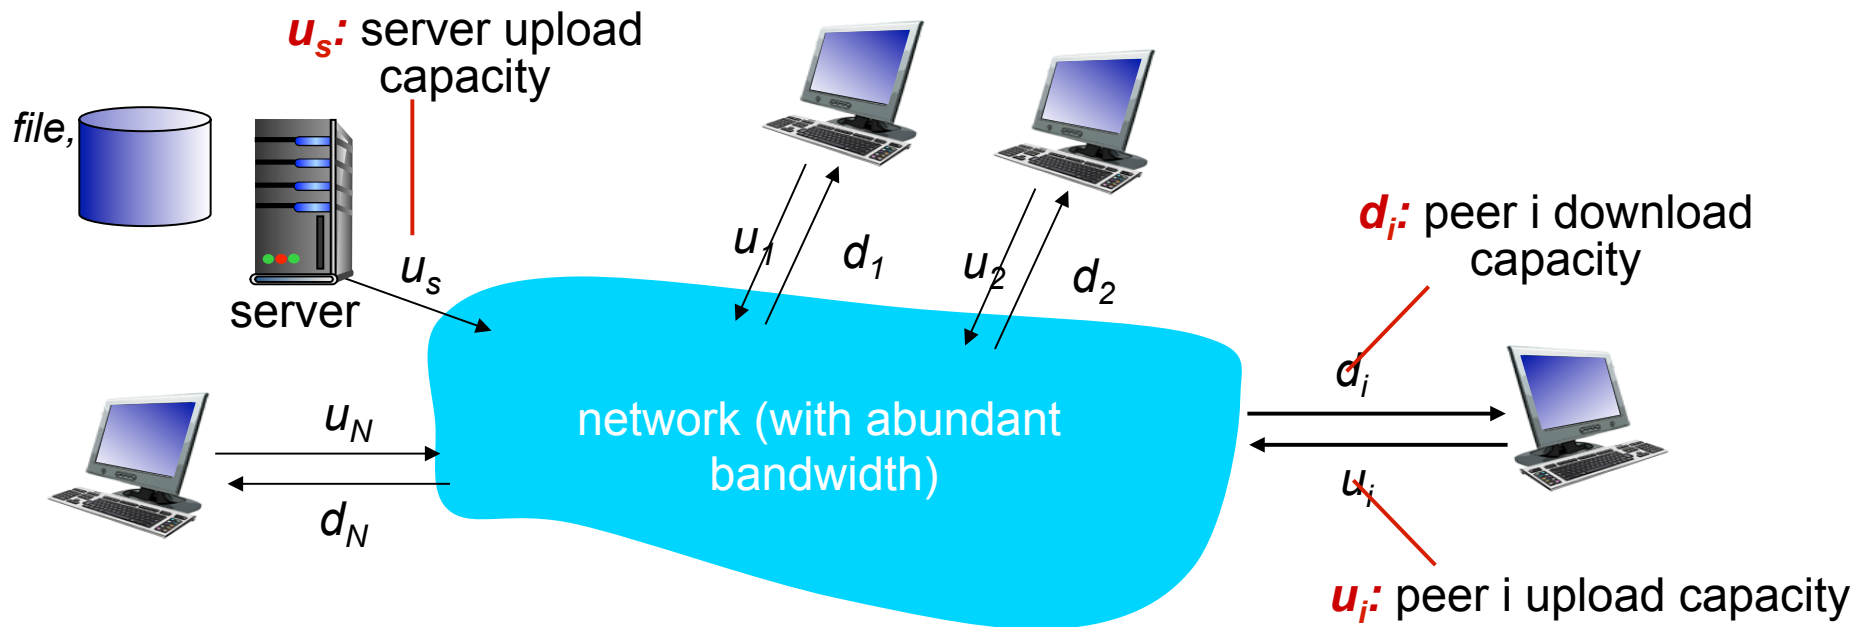
- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses
- Advantages
 - Distributes load of serving files.
- Challenges
 - How to find resources
 - Fairness
- *examples:*
 - file distribution (BitTorrent), Streaming (KanKan), VoIP (Skype)



File distribution: client-server vs P2P


Question: how much time to distribute file (size F) from one server to N peers?

– peer upload/download capacity is limited resource



File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:
 - time to send one copy: F/u_s
 - time to send N copies: NF/u_s
- ❖ **client:** each client must download file copy
 - d_{\min} = min client download rate
 - min client download time: F/d_{\min}




*time to distribute F
to N clients using
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N

File distribution time: P2P

- **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- ❖ **client:** each client must download file copy
 - min client download time: F/d_{\min}
- ❖ **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$




time to distribute F
to N clients using
P2P approach

$$D_{P2P} > \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

Comparing Client-server, P2P architectures




P2P file distribution: BitTorrent

- ❖ file divided into 256Kb chunks
- ❖ peers in torrent send/receive file chunks


tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file



P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
 - ❖ while downloading, peer uploads chunks to other peers
 - ❖ peer may change peers with whom it exchanges chunks
 - ❖ *churn*: peers may come and go
 - ❖ once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

requesting chunks:

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- ❖ Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

Food for thought...

- While P2P applications scale well to begin with, there are some concerns about the effect of mass-usage!
- On what assumptions is a packet-switching network based?
- On what assumptions is BitTorrent based?