


Software defined networking (SDN)

- Internet network layer: historically has been implemented via distributed, per-router approach
 - *monolithic* router contains switching hardware, runs proprietary implementation of Internet standard protocols (IP, RIP, IS-IS, OSPF, BGP) in proprietary router OS (e.g., Cisco IOS)
 - different “middleboxes” for different network layer functions: firewalls, load balancers, NAT boxes, ..
- ~2005: renewed interest in rethinking network control plane

Traffic engineering: difficult traditional routing




Q: what if network operator wants u-to-z traffic to flow along *uvwz*,
x-to-z traffic to flow *xwyz*?

A: need to define link weights so traffic routing algorithm computes
routes accordingly (or need a new routing algorithm)!

Link weights are only control “knobs”: wrong!


Traffic engineering: difficult



Q: what if network operator wants to split u-to-z traffic along uvwz *and* uxyz (load balancing)?

A: can't do it (or need a new routing algorithm)

Traffic engineering: difficult



Q: what if w wants to route blue and red traffic differently?

A: can't do it (with destination based forwarding, and LS, DV routing)

Software defined networking (SDN)

4. programmable control applications

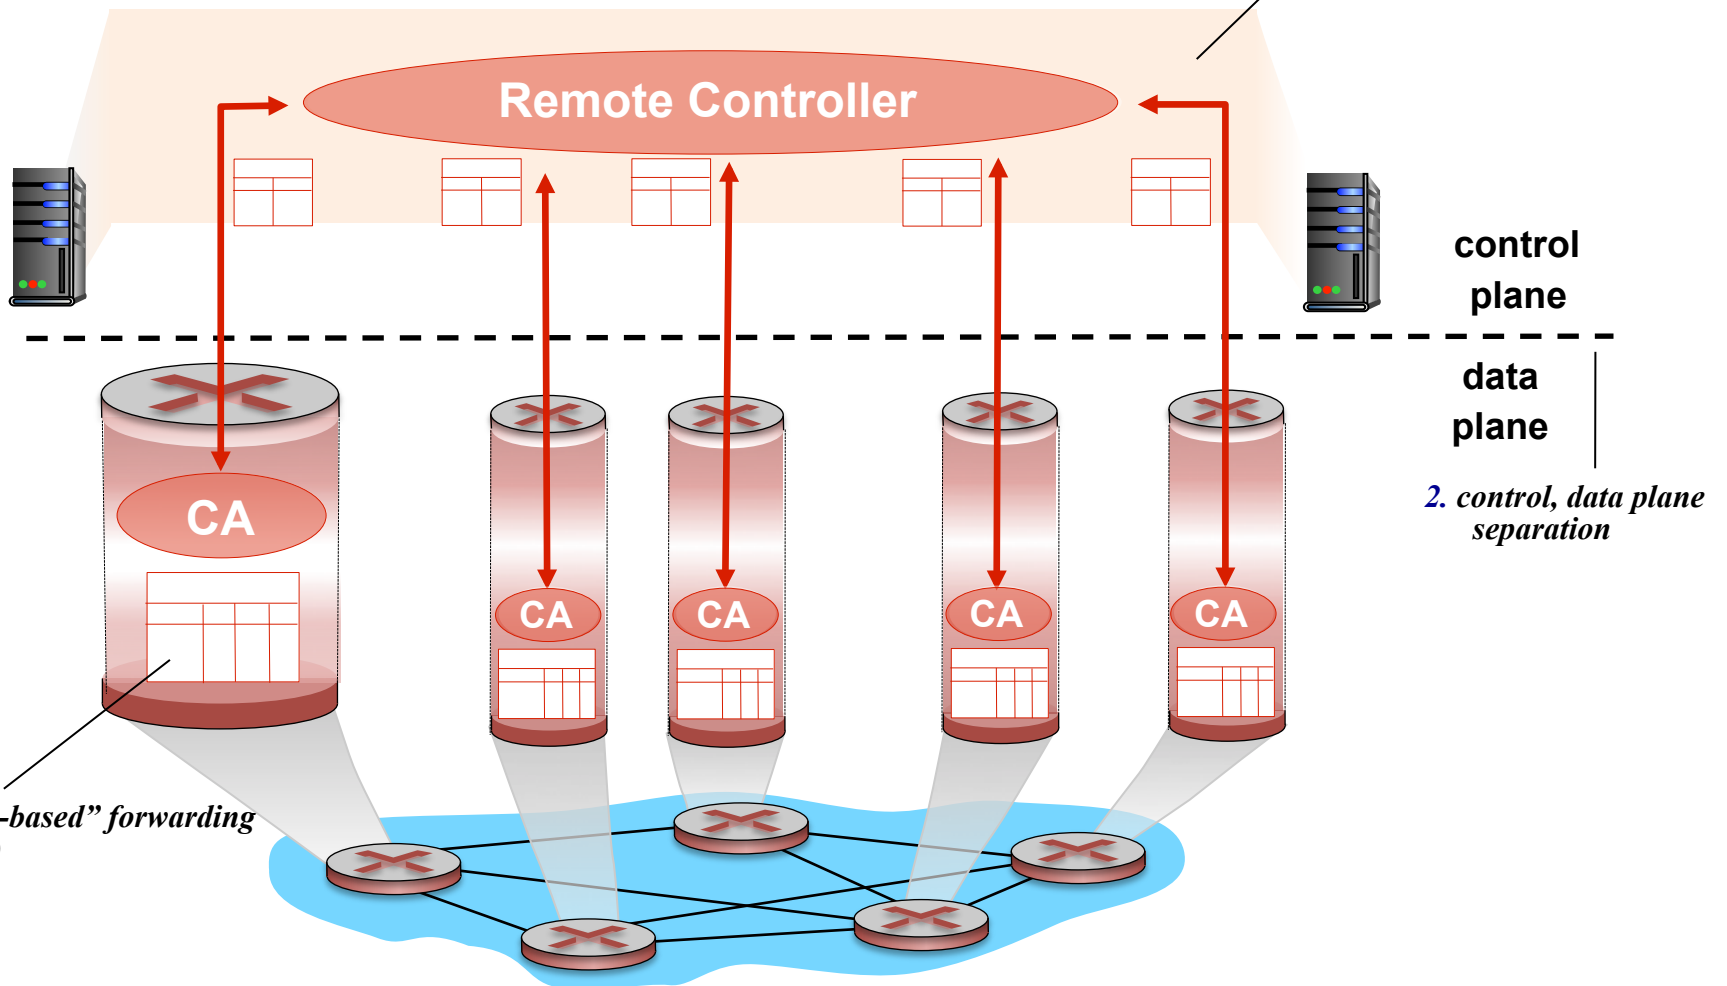
routing

access control

...

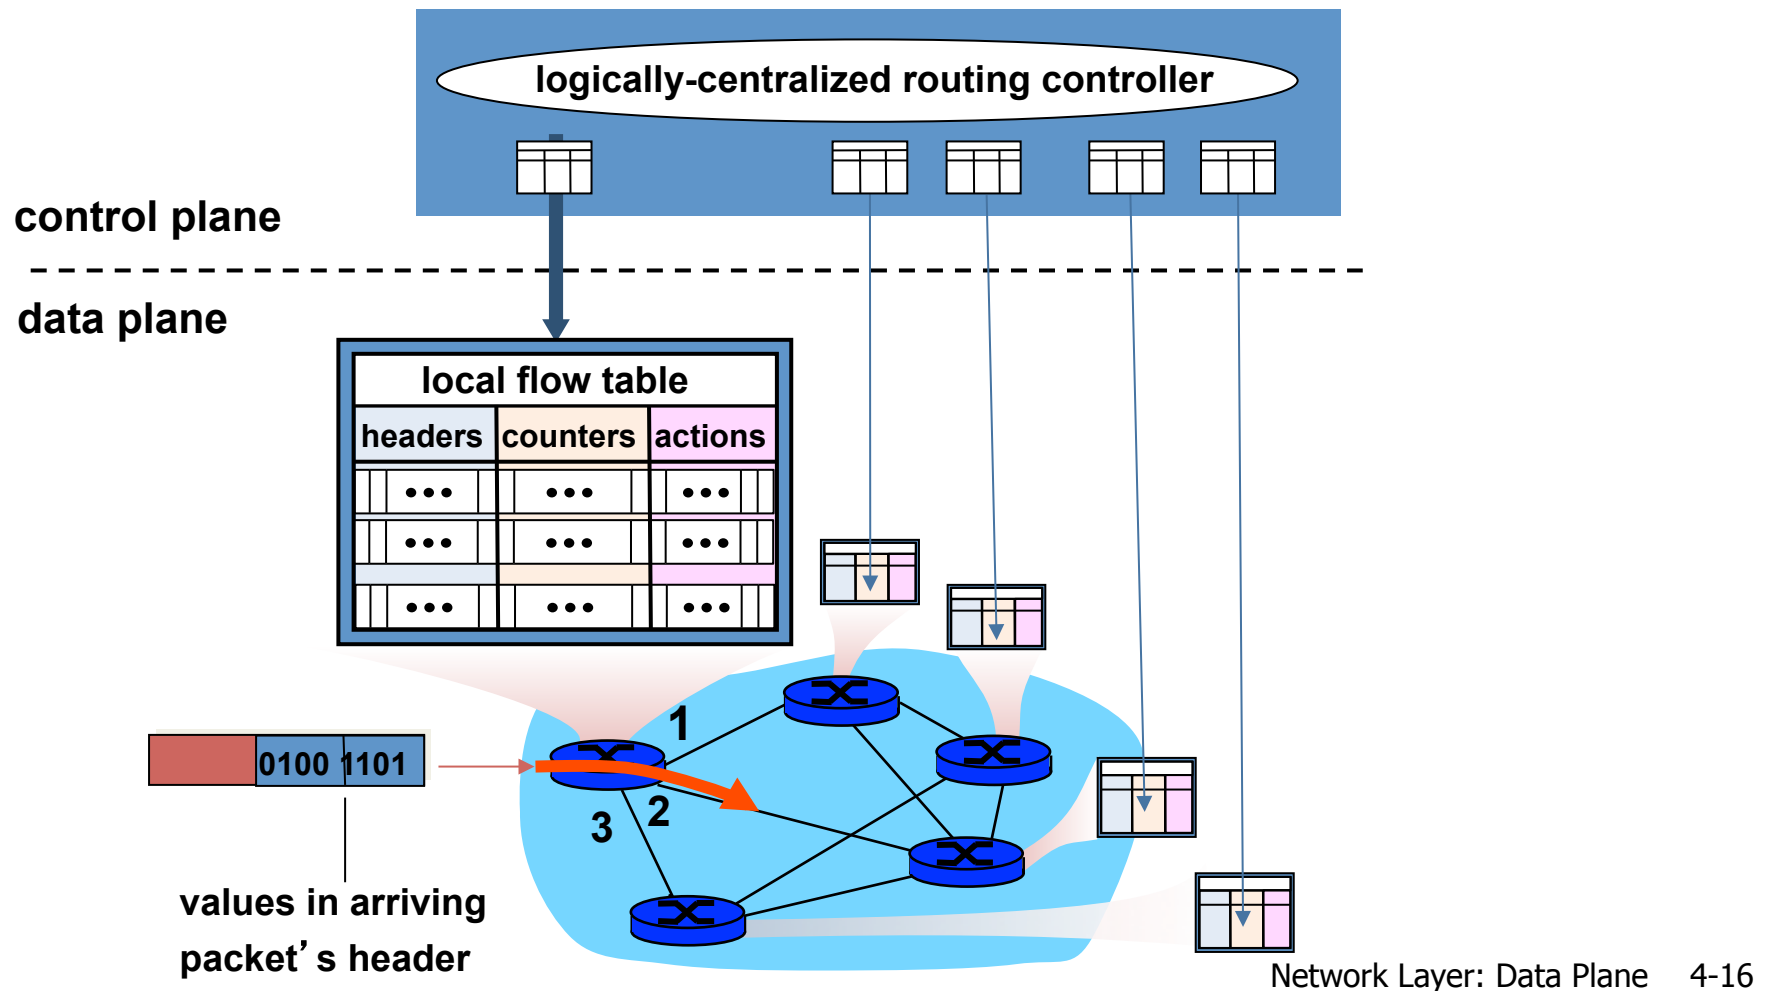
load balance

3. control plane functions external to data-plane switches

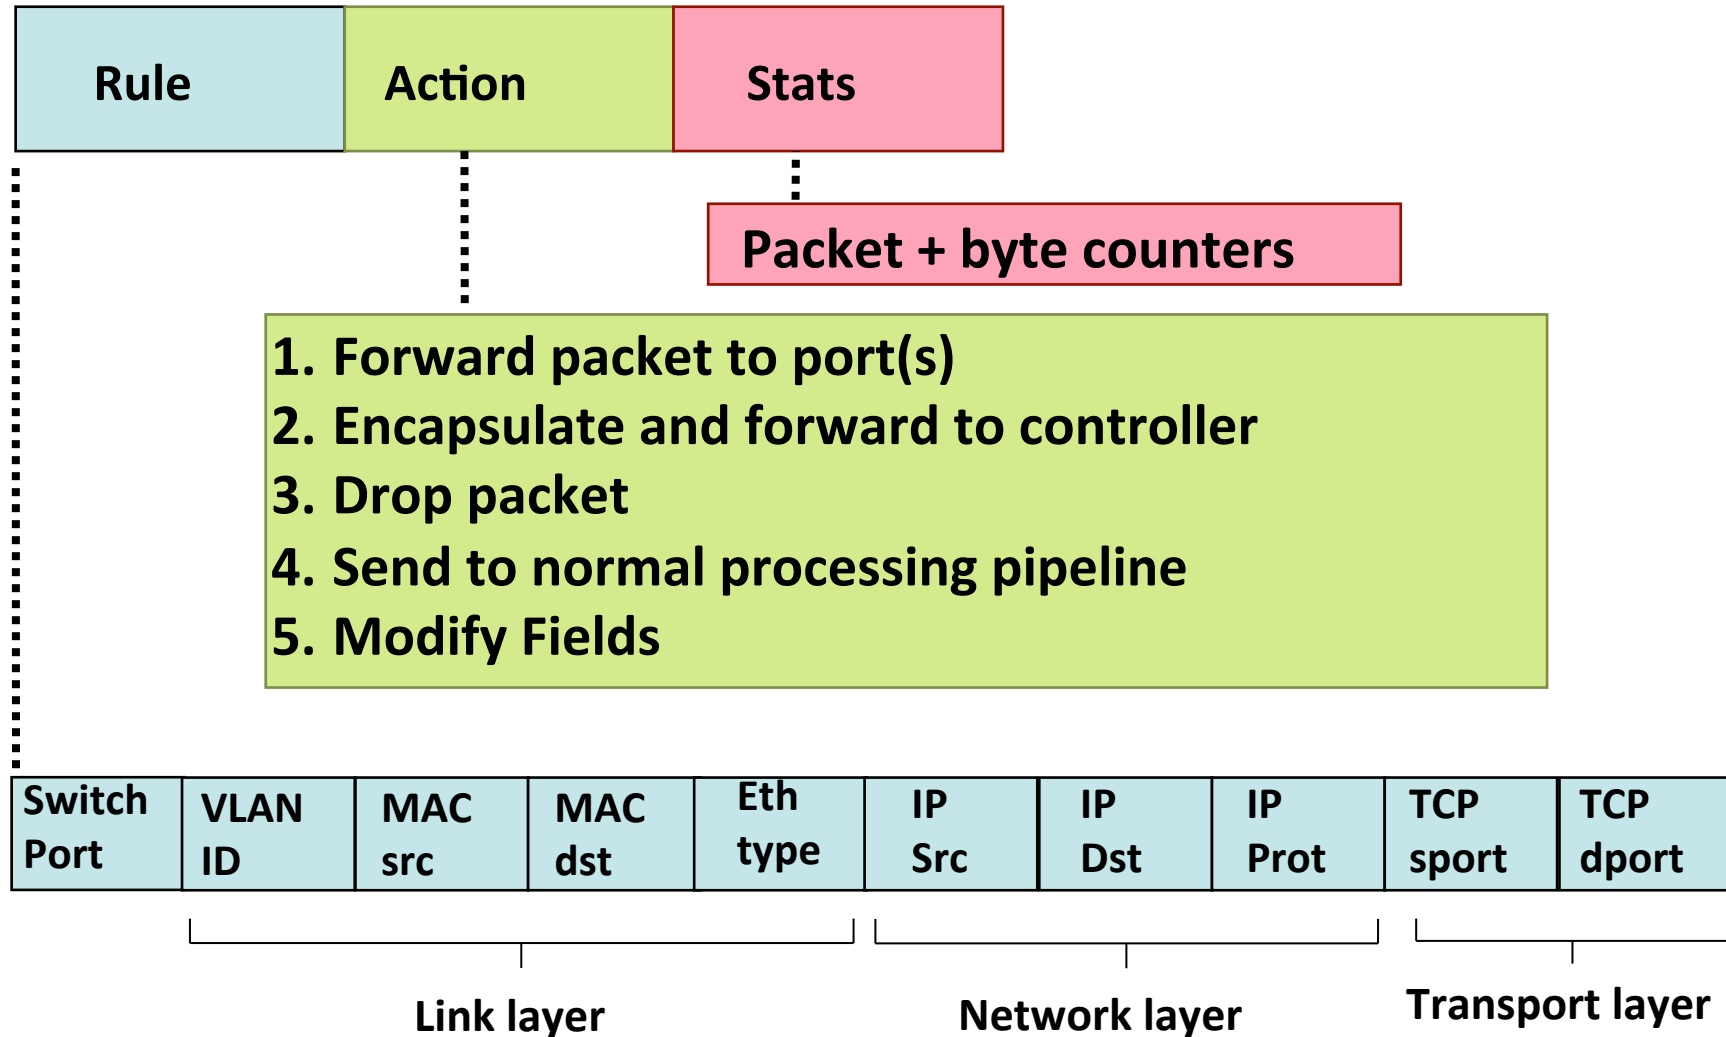


Generalized Forwarding and SDN

Each router contains a **flow table** that is computed and distributed by a **logically centralized routing controller**



OpenFlow: Flow Table Entries



Examples

Destination-based forwarding:

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|--------|
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|--------|

* * * * * 51.6.0.8 * * * port6

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

Firewall:

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Forward |
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|---------|
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|---------|

* * * * * * * * 22 drop

do not forward (block) all datagrams destined to TCP port 22

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Forward |
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|---------|
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|---------|

* * * * * 128.119.1.1 * * * drop

do not forward (block) all datagrams sent by host 128.119.1.1

IPv6

- IP was designed in the 1970's to support researchers and the military. The global use today wasn't predicted.
- Early design decisions led to predictions in the early 1990's of the "death of the Internet" due to
 - Shortage of Internet protocol (IP) addresses
 - Routing table explosion
 - Shortage of network numbers
- New capabilities needed:
 - Improved security (or even *some* security!)
 - Support for QoS (real-time services)
 - Better multicasting support
 - Mobile computing
- Two paths could be taken
 - Retrofit larger addresses and functionality onto IPv4
 - Develop a new version of IP

•Early proposals

- CLNP (Connectionless Network Protocol)
- SIP (Simple IP)
- Pip(Paul's Internet protocol)
- SIPP (SIP plus Pip) became IPv6 or IPng

Expanding the address space

- By 1996, all of the class A networks, 62% of class B and 37% of class C networks were allocated. Predicted to run out of addresses in 2008.
- The temporary solution is Classless InterDomain Routing (CIDR)
- With the push towards IP enabled devices this is likely to be only a temporary patch. IPv6 aims for a more permanent solution.
- How big is “big enough” for the address space?
 - Current 32 bit space allows *billions* of addresses.
 - Inefficiencies of allocation are almost inevitable.
 - May want to identify experimental networks in a different address spaces.
 - Can we predict future use? Are we setting ourselves up to need IPv7?
- Should the address be variable length?

IPv6 addressing

- 128 bits was a compromise
 - Fixed addresses are easier to manage and program.
- There is nearly 1 IP address for every molecule on the earth's surface.
 - Even with *inefficient* allocation there should be at least 1000 addresses per square metre.
- 128 bit addresses are written in hex:x:x:x:x:x:x:x:x
 - Each x is 16-bits = 4 hex digits
 - Leading zeros are not required
 - Sequence of zero fields given by "::"
- Examples:
 - 1080:0:0:0:8:800:200C:417A = 1080::8:800:200C:417A
 - FF01:0:0:0:0:0:0:43 = FF01::43
 - 0:0:0:0:0:0:0:1 = ::1

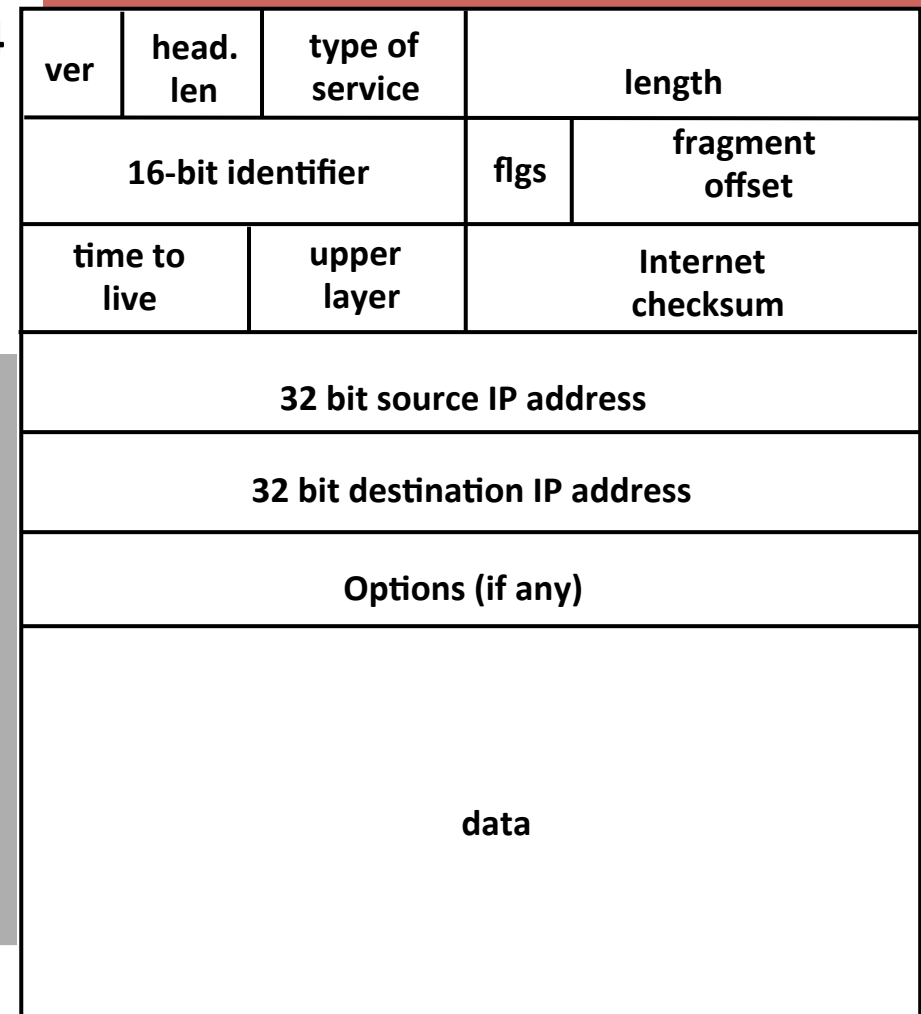
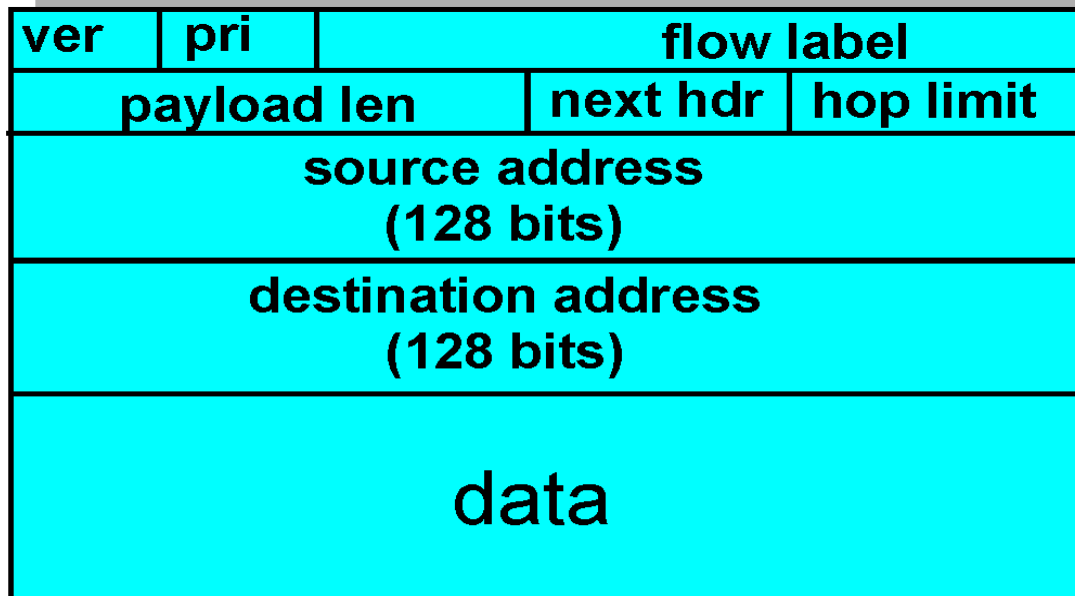
IPv6 Routing

- Each interface has an address (but possibly more than one):
 - Unicast – identify single interface
 - Anycast (or "cluster") – identify one out of number of interfaces ie. any one will do (eg. nearest host)
 - Multicast – identify set of interfaces, all of which are to receive message
- Intention in IPv6 addresses is to support ***hierarchical routing***
→ prefixes will **identify** registries, providers etc.
- Example: hierarchical organisation:

| | | | |
|-------------------|---------|-----------|--------------|
| s | n | m | 128-s-n-m |
| subscriber prefix | area id | subnet id | interface id |
- **This will help routing tables**
- Other possibilities:
 - global provider-based unicast address
 - geographic-based unicast address

IPv6 Header

- IPv6 8 fields in base header vs 13 fields in IPv4
- Faster processing
- Simpler management
- More flexibility



← 32 bits →

IPv4

Extension Headers

- This base header is followed by a number of optional ***extension headers***
 - Still allows flexibility
- Each header specifies code of next header/data component
- Extension headers commonly specify:
 - type of header and length
 - response if the router can't process the header — ignore the header, skip the packet, skip the packet and report an error

| Extension Header | Description |
|----------------------------|--------------------------------------|
| Hop-by-hop Options | Misc. information for routers |
| Routing | Full, or partial, route to follow |
| Fragmentation | Management of datagram fragments |
| Authorisation | Verification of Sender's Identity |
| Encrypted Security Payload | Information about encrypted Contents |
| Destination Options | Additional Info for destination |

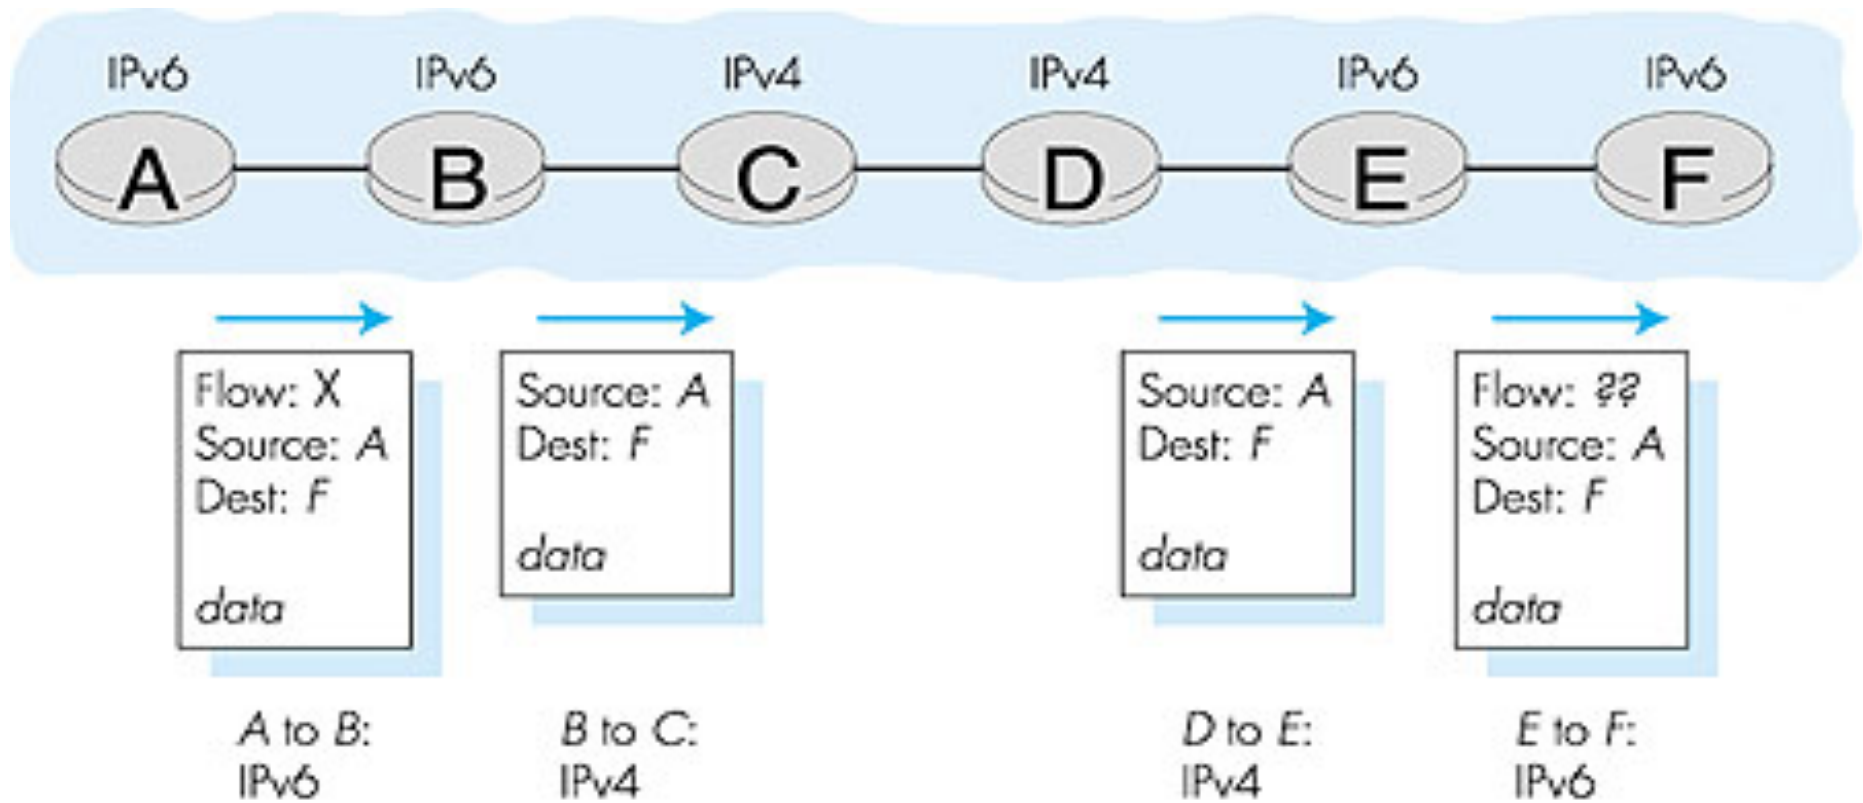
Fragmentation

- Fragmentation information is no longer in every header, but only in special extension headers
- Fragmentation is no longer performed at intermediate routers
 - The source host should choose datagram size so fragmentation is not necessary
 - Source host needs to run "path MTU discovery"
e.g. send sequence of datagram sizes to target until they don't arrive


Transition From IPv4 To IPv6

- Not all routers can be upgraded simultaneously
 - no “flag days”
 - How will the network operate with mixed IPv4 and IPv6 routers?
- Two proposed approaches:
 - *Dual Stack*: some routers with dual stack (v6, v4) can “translate” between formats
 - *Tunneling*: IPv6 carried as payload in IPv4 datagram among IPv4 routers
- Changes in networking technologies *always* proceed by evolution, rather than revolution

Dual Stack Approach



Tunneling



Tunneling

- Tunneling is a generically useful technique in networking...
 - Note that IPv6 payloads could be anything!
 - We could tunnel arbitrary protocols
- This is how experimentation is done –
 - New protocols are tunneled across existing infrastructure
 - Can spread new protocols across the network this way
 - Can link in “new” devices or “old” devices this way
 - Can implement secure networks this way (see later...)

IPv6 Summary

- What has changed
 - Simpler, fixed length header
 - No fragmentation in routers
 - Options in extension headers
 - No checksum
 - 128 bit (IPv4 32 bit) addresses, *with hierarchy*
 - Additional support for
 - Multicast and anycast routing
 - Security
 - Mobile hosts and autoconfiguration
 - Real Time applications